

```

#https://colab.research.google.com/drive/1EPraBTVrxihhg7T9OLIZdMcerg-FUUi_?usp=sharing
!pip install --quiet transformers torch

from transformers import pipeline
from IPython.display import display, Markdown

gpt2 = pipeline("text-generation", model="gpt2")
distilgpt2 = pipeline("text-generation", model="distilgpt2")
gpt2_medium = pipeline("text-generation", model="gpt2-medium")

prompt = "Write a haiku about AI and nature."

competitors = [
    "GPT2",
    "DistilGPT2",
    "GPT2-Medium"
]
answers = [
    gpt2(prompt, max_length=60, temperature=0.7)[0]["generated_text"],
    distilgpt2(prompt, max_length=60, temperature=0.7)[0]["generated_text"],
    gpt2_medium(prompt, max_length=60, temperature=0.7)[0]["generated_text"]
]

display(Markdown("## 🚀 Model Outputs\n"))
for comp, ans in zip(competitors, answers):
    display(Markdown(f"### 📄 {comp}\n{ans}\n---"))

def score_response(text):
    score = 0
    score += len(text.split())
    score += text.count("\n") * 5
    return score

scores = [(comp, score_response(ans)) for comp, ans in zip(competitors, answers)]
scores.sort(key=lambda x: x[1], reverse=True)

display(Markdown("## 🏆 Final Ranking (Best → Worst)\n"))
for i, (model, score) in enumerate(scores):
    display(Markdown(f"**Rank {i+1}: {model}** - Score: {score}"))

```



```
/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
warnings.warn(
Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster d
WARNING:huggingface_hub.utils._http:Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to en
config.json: 100%                               665/665 [00:00<00:00, 18.0kB/s]

model.safetensors: 100%                         548M/548M [00:06<00:00, 101MB/s]

Loading weights: 100%                           148/148 [00:00<00:00, 233.64it/s, Materializing param=transformer.wte.weight]
GPT2LMHeadModel LOAD REPORT from: gpt2
Key | Status | |
-----+-----+---
h.{0...11}.attn.bias | UNEXPECTED | |

Notes:
- UNEXPECTED :can be ignored when loading from different task/architecture; not ok if you expect identical arch.
generation_config.json: 100%                     124/124 [00:00<00:00, 4.54kB/s]

tokenizer_config.json: 100%                      26.0/26.0 [00:00<00:00, 579B/s]

vocab.json: 1.04M/? [00:00<00:00, 2.87MB/s]
merges.txt: 456k/? [00:00<00:00, 1.23MB/s]
tokenizer.json: 1.36M/? [00:00<00:00, 2.48MB/s]

config.json: 100%                               762/762 [00:00<00:00, 22.1kB/s]

model.safetensors: 100%                         353M/353M [00:06<00:00, 79.9MB/s]

Loading weights: 100%                           76/76 [00:00<00:00, 364.12it/s, Materializing param=transformer.wte.weight]
GPT2LMHeadModel LOAD REPORT from: distilgpt2
Key | Status | |
-----+-----+---
transformer.h.{0, 1, 2, 3, 4, 5}.attn.bias | UNEXPECTED | |

Notes:
- UNEXPECTED :can be ignored when loading from different task/architecture; not ok if you expect identical arch.
generation_config.json: 100%                     124/124 [00:00<00:00, 6.19kB/s]

tokenizer_config.json: 100%                      26.0/26.0 [00:00<00:00, 305B/s]

vocab.json: 1.04M/? [00:00<00:00, 2.47MB/s]
merges.txt: 456k/? [00:00<00:00, 3.16MB/s]
tokenizer.json: 1.36M/? [00:00<00:00, 1.84MB/s]

config.json: 100%                               718/718 [00:00<00:00, 63.7kB/s]
```